
Large-Scale Distributed Systems

Project 3: Firefly-inspired synchronization

Laurent HAYEZ

December 14, 2015

Table of contents

1	Introduction	1
2	Protocol skeleton	1
3	Phase-advance and phase-delay protocols	2
3.1	Implementation	2
3.2	Analysis of the protocols	2
4	Conclusion	6

1 Introduction

In nature, fireflies produce light in order to attract mates or prey. One interesting feature of these beetles is that when they emit light in group, at some point, they do it in a synchronized manner, just by looking at when their neighbours emit light. This feature is interesting in large-scale distributed systems, as synchronization might be required, but one node does not know every other nodes.

The objective is thus to inspire ourselves from fireflies to try to synchronize nodes in a decentralized manner. At first we will detail the protocol skeleton and explain how the core of the protocols will work. Then we will look at a two models called “phase-advance” and “phase-delay” and briefly analyze them. The main and final part will be the “adaptive Ermentrout model” which is more representative of the reality. We will explain the implementation specificities and analyze this model in different situations.

2 Protocol skeleton

According to the paper “Firefly-inspired Heartbeat Synchronization in Overlay Networks”, the skeleton for the different algorithms is composed of two main functions, namely `ACTIVE_THREAD` and `PASSIVE_THREAD`. We provide the pseudo code for the implementation in Algorithm 2.1.

In the different protocols, a node is an oscillator characterized by its phase φ and the cycle length Δ . We define φ as a sawtooth function with domain $[0, 1]$ such that $\frac{d\varphi}{dt} = \frac{1}{\Delta}$. This is represented in Figure 1.

When φ reaches 1, the node will send a flash to a set of neighbour nodes, and φ is reset to 0. The cycle length, depending on the model chosen, can be the same or different for all nodes. The function `UPDATEPHI` will differ in our implementations, but we will come back on this when needed.

The core of the synchronization protocol is the function `PROCESSFLASH`, i.e. what a node does when it receives a flash. This function is responsible of how φ is updated. Depending on the implementation, φ will be updated, or Δ will be updated and will affect φ .

The underlying overlay network protocol used for a node to know its neighbours is the Peer Sampling Service (PSS).




Figure 1: Representation of φ and its relation with Δ

We briefly discuss how the protocol skeleton works before we move to the implementation of the phase-advance and phase-delay protocols.

For every implementations, we start by initializing the PSS at each node, and wait two minutes to have a consistent view. Then we start a thread `ACTIVETHREAD` and a periodic thread `UPDATEPHI`. As we defined φ to be such that $\frac{d\varphi}{dt} = \frac{1}{\Delta}$, we have $\varphi = \int \varphi dt$, which explains how we update φ . If $\varphi \geq 1$, we fire the event “Flash!” that triggers `ACTIVETHREAD` to send that it emitted a flash to the nodes in its view. In `UPDATEPHI`, we then set φ to 0, as in Figure 1 and call `ACTIVETHREAD` to wait for a new flash.

3 Phase-advance and phase-delay protocols

3.1 Implementation

The fundamental difference with the protocol skeleton here is the implementation of the `PROCESSFLASH` function. As the two protocols are almost the same, we added the boolean `PHASE_ADVANCE` to signify the usage of the phase-advance or phase-delay protocol. The pseudo-code of the implementation is given in Algorithm 3.1.

3.2 Analysis of the protocols

We tried our implementation with 100 nodes per experiment and $\Delta = 1, 2$ and 5. For the PSS, we used a view of 10 peers with a period of 20 seconds for each update. We used a random selection for the peer selection, and we used a swapping parameter of 3 and a healing parameter of 2 for the view exchange policy.

Algorithm 2.1 Skeleton for the Firefly algorithms

Variables:

φ \triangleright phase
 Δ \triangleright cycle length

$$\text{update_phi_period} = \begin{cases} \frac{\Delta}{10} & \text{if } \Delta < 1 \\ \frac{1}{10\Delta} & \text{if } \Delta \geq 1 \end{cases}$$

function SENDFLASH()

$P \leftarrow$ view from PSS

 send flash to all peers in P

end function

function PROCESSFLASH()

 depends on the implementation

end function

function UPDATEPHI()

if $\varphi < 1$ **then**

$\varphi \leftarrow \varphi + \frac{1}{\Delta} \cdot \text{update_phi_period}$

else

 fire event “Flash!”

$\varphi \leftarrow 0$

 start new thread ACTIVETHREAD

end if

end function

function ACTIVETHREAD()

 wait for the event “Flash!”

 sendFlash()

end function

function PASSIVETHREAD()

 receive flash

 processFlash()

end function

Algorithm 3.1 processFlash for the phase-advance and phase-delay protocols

Variables:

phase_advance \leftarrow true or false

function PROCESSFLASH()

if phase_advance **then**

$\varphi \leftarrow 1$

else

$\varphi \leftarrow 0$

end if

end function

We start by analyzing the convergence of the phase-advance protocol. On Figure 2, we can see that from time to time, the protocol seems to synchronize, it does not converge to a stable state where all the nodes emit flashes at the same time. In fact, what happens is what we see on Figures 3 and 4. Whenever a node flashes, it tells ten other nodes, that set φ to 1, emit a flash, tell ten other nodes, which emit a flash and so on. We thus have an exponential growth of the flash rate which makes the memory used explode very fast. In the logs, we had this very error:

```
17:37:43.276437 (66) Too much memory used:  wants 2097182 (max:
2097152), end of process.
```

This model is therefore highly impractical because it overloads the network very quickly. By analogy with fireflies, at first a few fireflies would emit light, and then by observing the neighbours they would emit light every time they see a neighbour emitting light, quickly resulting in emitting light constantly, which does not happen in reality.




Figure 2: Phase-advance model synchronization with $\Delta = 1$




Figure 3: Phase-advance model synchronization with $\Delta = 2$

Now let's take a look at the phase-delay model. On Figures 5 and 6, we represent the emissions for $\Delta = 1$ and $\Delta = 2$ respectively. On both figures we cannot distinguish

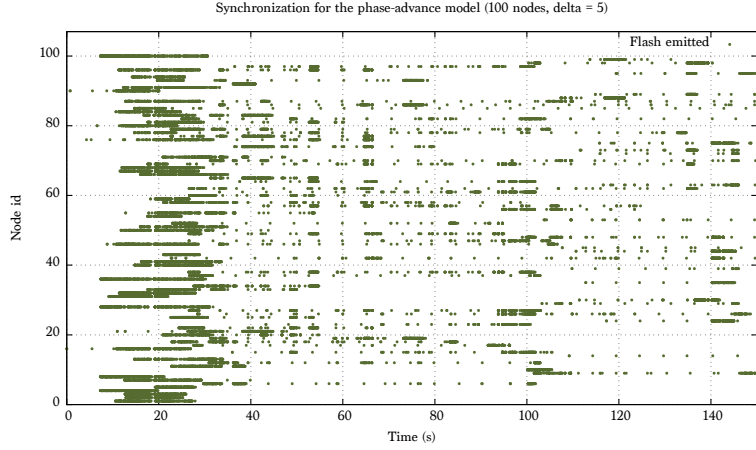


Figure 4: Phase-advance model synchronization with $\Delta = 5$

a convergence in the flash emissions. On Figure 7 however, we can see that from 150 seconds after the first flash, the nodes that emit flashes are synchronized.

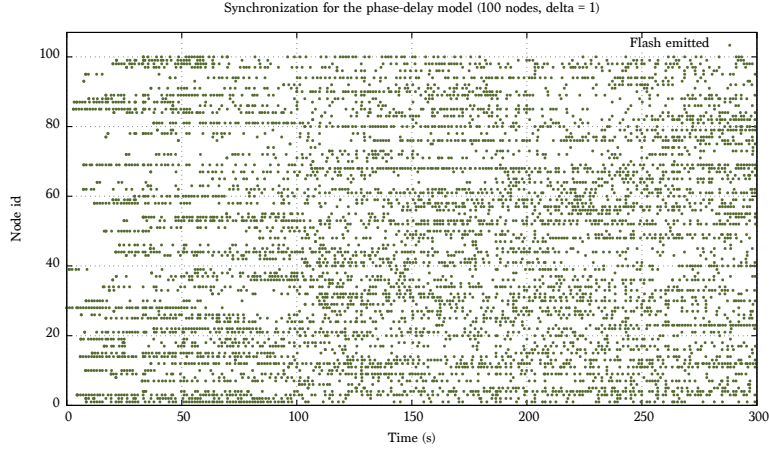


Figure 5: Phase-delay model synchronization with $\Delta = 1$

Although we seem to obtain some synchronization for $\Delta = 5$, the phase-delay protocol is not guaranteed to converge to a stable state where all the nodes are synchronized. Both phase-advance and phase-delay protocols rely on the assumption that when a message is sent, it is instantly received. In practice, this assumption is highly unrealistic due to packet loss, congestion or simply because even the fastest networks, a message is not instantly delivered.

In our experiments, we considered Δ to be fixed and the same at all nodes. This assumption does not represent reality either; why would all fireflies emit light at the same rate? If we had different Δ s at all nodes, or if the skew of the clocks is important, none of the models is guaranteed to converge, and it is rather logical as we don't update Δ . If a node starts with $\Delta = 4.5$ and another with $\Delta = 5.5$, they will never synchronize. Therefore, both phase-advance and phase-delay are mathematically insufficient to model the reality.




Figure 6: Phase-delay model synchronization with $\Delta = 2$

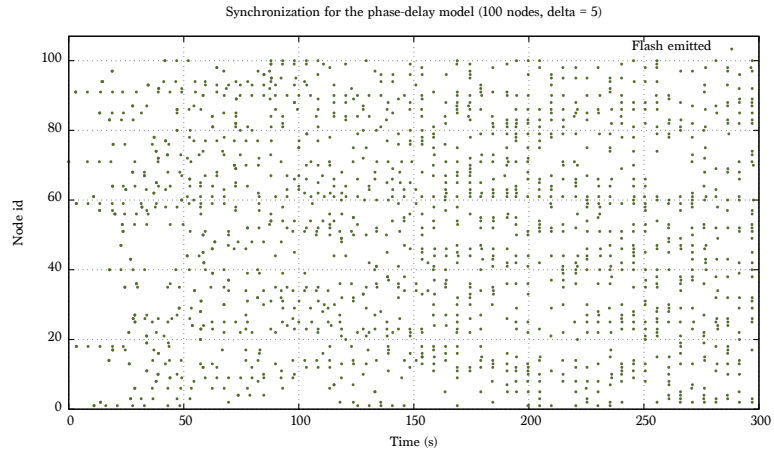


Figure 7: Phase-delay model synchronization with $\Delta = 5$

4 Conclusion